


Discussions

This tutorial discussion questions are based on the following pizza database schema. The ER diagram is shown below.



The ER diagram produces the following schemas:

Relation	Description
Pizzas(<u>pizza</u>)	All the pizzas of interest.
Customers(<u>cname</u> , area)	The name and location of each customer.
Restaurants(<u>rname</u> , area)	The name and location of each restaurant.
Recipes(<u>pizza</u> , <u>ingredients</u>)	The ingredients used in each pizza.
Sells(<u>rname</u> , <u>pizza</u> , price)	Pizzas sold by restaurants and the prices.
Likes(<u>cname</u> , <u>pizza</u>)	Pizzas that customers like.

Additionally, we have the following foreign key constraints on the database schema:

- $(Recipes.pizza) \rightsquigarrow (Pizzas.pizza)$
- $(Sells.rname) \rightsquigarrow (Restaurants.rname)$
- $(Sells.pizza) \rightsquigarrow (Pizzas.pizza)$
- $(Likes.cname) \rightsquigarrow (Customers.cname)$
- $(Likes.pizza) \rightsquigarrow (Pizzas.pizza)$

1. (Basic Query)

- (a) Find all restaurant names located in the 'East'.
- (b) Find all areas with at least one customer.
- (c) Find all the pizzas sold by 'Corleone Corner'.

2. (Equivalent Query) For each of the following queries, write an *equivalent* SQL query that does **not** use any subquery. You may use set operation.

(a) Query A

```
1  SELECT DISTINCT cname
2  FROM    Likes L
3  WHERE   EXISTS (
4      SELECT 1
5      FROM    Sells S
6      WHERE   S.rname = 'Corleone Corner' AND S.pizza = L.pizza
7  );
```

(b) Query B

```
1  SELECT cname
2  FROM    Customers C
3  WHERE   NOT EXISTS (
4      SELECT 1
5      FROM    Likes L, Sells S
6      WHERE   S.rname = 'Corleone Corner'
7          AND S.pizza = L.pizza AND C.cname = L.cname
8  );
```

(c) Query C

```
1  SELECT DISTINCT rname
2  FROM    Sells
3  WHERE   rname <> 'Corleone Corner'
4      AND price > ANY (
5      SELECT price
6      FROM    Sells
7      WHERE   rname = 'Corleone Corner'
8  );
```

(d) Query D

```
1  SELECT rname, pizza, price
2  FROM    Sells S
3  WHERE   price >= ALL (
4      SELECT S2.price
5      FROM    Sells S2
6      WHERE   S2.rname = S.rname
7          AND S2.price IS NOT NULL
8  );
```

-
3. **(SQL Query)** Write an SQL query to answer each of the following questions on the pizza database *without using aggregate functions*. Remove duplicate records from all query results.
- (a) Find pizzas that Moe likes but is not liked by Lisa.
 - (b) Find pizzas that are sold by at most one restaurant in each area; exclude pizzas that are not sold by any restaurant.
 - (c) Find all tuples (A, P, P_{min}) where P is a pizza that is available in area A (*i.e.*, there is some restaurant in area A selling pizza P) and P_{min} is the *lowest* price of P in area A .
 - (d) Find the most expensive pizzas and the restaurants that sell them (*at the most expensive price*).

Challenge

The answers to the following questions is given without explanation. Please discuss them on Canvas.

1. **(SQL Query)** Write an SQL query to answer each of the following questions on the pizza database *without using aggregate functions*. Remove duplicate records from all query results.
 - (a) Find all tuples (A, P, P_{min}, P_{max}) where P is a pizza that is available in area A (*i.e.*, there is some restaurant in area A selling pizza P), P_{min} is the *lowest* price of P in area A and P_{max} is the *highest* price of P in area A .
2. **(Update)** Consider the following relational schema.

```
1 CREATE TABLE Offices (  
2     office_id      INT,  
3     building       TEXT NOT NULL,  
4     level          INT  NOT NULL,  
5     room_number    INT  NOT NULL,  
6     area           INT,  
7     PRIMARY KEY   (office_id),  
8     UNIQUE         (building, level, room_number)  
9 );
```

```
1 CREATE TABLE Employees (  
2     emp_id         INT,  
3     name           TEXT NOT NULL,  
4     office_id      INT  NOT NULL,  
5     manager_id     INT,  
6     PRIMARY KEY    (emp_id),  
7     FOREIGN KEY    (office_id) REFERENCES Offices (office_id)  
8         ON UPDATE CASCADE,  
9     FOREIGN KEY    (manager_id) REFERENCES Employees (emp_id)  
10        ON UPDATE CASCADE  
11 );
```

Suppose that the office with `office_id = 123` needs to be renovated. Write an SQL statement to reassign the employees located in this office to another temporary office located at room number 11 on level 5 at the building named *Tower1*.

Hint: You can use subquery in an update statement.

3. **(Backward Reasoning)** You are given the following schema:
 - Students(matric, sname)
 - Workings(pid, matric, since)
 - Projects(pid, pname)
 - Categories(pid, cname)

You are also given the following foreign key:

- $(Workings.matric) \rightsquigarrow (Students.matric)$
- $(Workings.pid) \rightsquigarrow (Projects.pid)$
- $(Categories.pid) \rightsquigarrow (Projects.pid)$

Consider the following SQL query:

```
1 SELECT *
2 FROM   Students NATURAL JOIN Workings
3         NATURAL JOIN Projects
4         NATURAL JOIN Categories;
```

Say it produces the following result:

matric	sname	pid	since	cname
A0001	AA	P01	2002	CA
A0001	AA	P01	2002	CB
A0001	AA	P02	2004	CB
A0002	BB	P01	2003	CA
A0002	BB	P01	2003	CB
A0003	CC	P03	2004	CA
A0003	CC	P03	2004	CC
A0003	CC	P03	2004	CD
A0004	AA	P03	2004	CA
A0004	AA	P03	2004	CC
A0004	AA	P03	2004	CD

Now consider another the query to find all pair of distinct projects' pid (*i.e.*, $(p1, p2)$) such that the two projects have exactly the same set of categories¹. It produces the following result:

pid	pid
P01	P04
P04	P01
P03	P05
P05	P03

¹This involves some constructs that you will only learn in Lecture 06.

Simply note that P01 and P04 have exactly the same set of categories. Similarly, P03 and P05 have exactly the same set of categories. What is a possible result of the following SQL query? Show your answer using only P01, P03, P04, and P05?

```
1 SELECT pid FROM Categories
2 EXCEPT ALL
3 SELECT DISTINCT pid FROM Categories WHERE cname <> 'CA';
```