

Discussions

1. (**Superkey**) Consider the following relation instance r of the relational schema $R(A, B, C, D)$.

R			
A	B	C	D
0	0	0	1
2	1	2	0
1	1	2	0
0	0	1	2

- (a) Assuming that r is a *valid* relation instance of R , write down all the *possible* superkeys of R .
- (b) Additionally, suppose that it is also known that $\{A, C\}$ is *definitely* a superkey of R . Based on the additional information, write down all the *possible* candidate keys of R .
2. (**SQL DDL**) Consider a relational database for a company that consists of the following two tables.
- `Offices(office_id, building, floor, room_number, area)`
 - `Employees(emp_id, name, office_id, manager_id)`

The database satisfies the following constraints:

1. **Offices** stores information about the office rooms in the company. Each room has a unique identifier `office_id` (which is the primary key of **Offices**) and information on its building name, floor level, room number and floor area.
2. $\{\text{building, floor, room_number}\}$ is a *candidate key* of **Offices**.
3. **Employees** stores information about the employees in the company.
4. `emp_id` is the *primary key* of **Employees**.
5. The name of each employee must be a non-**NULL** value.
6. Each employee must be assigned to exactly one office identified by `office_id`.
7. Each employee may be managed by at most one manager.
8. If an employee is managed by someone, the `emp_id` of his/her manager is recorded in `manager_id`.
9. A record in **Offices** cannot be removed if there's some employee assigned to that office.
10. A manager in **Employees** cannot be removed if there's some other employee managed by that manager.
11. Any modification to `office_id` in **Offices** is propagated to other database records.
12. Any modification to `emp_id` in **Employees** is propagated to other database records.

Write SQL statements to create the database schema with appropriate attribute domains and constraints.

3. (SQL DDL) Consider a relational database for an online shop that consists of the following five tables.

- Books(isbn, title, authors, year, edition, publisher, number_pages, price)
- Customers(cust_id, name, email)
- Carts(cust_id, isbn)
- Purchase(pid, purchase_date, cust_id)
- Purchased_items(pid, isbn)

The database satisfies the following constraints:

1. **Books** records information about the books available for sale in an online shop. Each book has a unique identifier **isbn** and information about its title, authors, publishers, publication year, edition, number of pages and selling price. The title and authors must have non-NULL values. The value of the edition must be non-NULL with one of the following values: *paperback*, *hardcover* or *ebook*. The selling price must have a positive value. If the number of pages is known, it must be a positive value.
2. **Customers** stores information about the shop's customers. Each customer has a unique identifier **cust_id**, a name and an email address. The name must have a non-NULL value.
3. **Carts** stores information about the books in customers' shopping carts. Each shopping cart record indicates a book that a customer is interested to purchase but has not yet purchased.
4. When a customer decides to purchase their selected books, a new record for this purchase is recorded in the **Purchase** table which has the following information:
 - (a) A unique identifier **pid**.
 - (b) The date of the purchase.
 - (c) The customer identifier.In addition, each book in the customers' shopping cart is added to the **Purchased_items** table and the customers' shopping cart is emptied.

- (a) Write SQL statements to create a database schema with appropriate attribute domains and constraints.
- (b) Suppose that the schema of **Purchase**(pid, purchase_date, cust_id) is changed such that the attribute **purchase_date** being replaced by the attribute **purchase_timestamp**, where each customer has at most one purchase at any timestamp¹. How would this change affect your answer for part a?
- (c) For each of the following additional constraints on the database, state whether the constraint can be expressed using SQL constructs that you have learned. If it is possible, add this constraint to your database schema in part a.

1. If a book is a hardcover edition, its selling price must be at least 30.
2. If a book has both hardcover and paperback editions (for the same book title and authors), the selling price for the hardcover edition must be higher than the selling price for the paperback edition.
3. If the number of pages in a book is more than 1000, the edition of the book must be an ebook or its price must be at least 100.
4. All the books published by "Acme" from 2010 onwards have only ebook edition.

¹<https://www.postgresql.org/docs/15/datatype-datetime.html>

Challenge

The answers to the following questions is given without explanation. Please discuss them on Canvas.

1. **(SQL DDL)** This question refers to your solution for question [3](#).

(a) Consider the following additional constraints on the database:

1. If a customer is deleted from **Customers**, remove all the customers' records from **Carts** and **Purchase**.
2. If a book is deleted from **Books**, remove all records from **Carts** that reference this book. Additionally, for each of the records in **Purchased_items** that reference this book, change its **isbn** value to the default value of 0.
3. If a purchase is deleted from **Purchase**, remove all the records from **Purchased_items** that reference this purchase.
4. Any modification of **cust_id** in **Customers** is propagated to other records in the database.
5. Any modification of **isbn** in **Books** is propagated to other records in the database.
6. Any modification of **pid** in **Purchase** is propagated to other records in the database.

Add these additional constraints to your answer for question [3a](#).