

CS2106 Introduction to Operating Systems

Semester 2 2023/2024

Questions

Tutorial 3: Process Scheduling

1. [Understanding scheduling algorithm] (Adapted from Andrew.T) Six batch jobs. *A* through *F*, arrive at a computer center at almost the same time. The estimated running time and priorities are as follows:

Process Name	Running time (ms)	Priority (1 is highest priority)
A	8	4
B	5	6
C	4	3
D	6	2
E	1	5
F	2	1

For each of the following scheduling algorithms, determine the mean process turnaround time. Ignore process switching overhead. Assume all processes are **CPU-bound** and that only one job at a time runs, until it finishes (i.e., non-preemptive).

- (a) Priority scheduling.
- (b) First Come First Serve (run in the order A through F)
- (c) Shortest Job First

-
2. [Walking through Scheduling Algorithms] Consider the following execution scenario:

Program A, Arrives at time 0
Behavior (CX = Compute for X Time Units, IOX = I/O for X Time Units): C3, IO1, C3, IO1

Program B, Arrives at time 0
Behavior: C1, IO2, C1, IO2, C2, IO1

Program C, Arrives at time 3
Behavior:

- a) Show the scheduling time chart with First-Come-First-Serve algorithm. For simplicity, we assume all tasks block on the **same I/O resource**.

Below is a sample sketch up to time 1:



- b) What are the turnaround time and the waiting time for program A, B and C? In this case, waiting time includes all time units where the program is ready for execution but could not get the CPU.
- c) Use the Round **Robin** algorithm to schedule the same set of tasks. Assume a time quantum of **2 time units**. Draw out the scheduling time chart. State any assumptions you may have.
- d) What is the response time for tasks A, B and C? In this case, we define response time as the time difference between the arrival time and the first time when the task receives CPU time.

3. [MLFQ] As discussed in the lecture, the simple MLFQ has a few shortcomings. Describe the scheduling behavior for the following two cases.

1. (Change of heart) A process with a lengthy CPU-intensive phase followed by I/O-intensive phase.
2. (Gaming the system) A process repeatedly gives up CPU just before the time quantum lapses.

The following are two simple tweaks. For each of the rules, identify which case (1 or 2 above) it is designed to solve, then briefly describe the new scheduling behavior.

- i. (Rule – Accounting matters) The CPU usage of a process is now accumulated across time quanta. Once the CPU usage exceeds a single time quantum, the priority of the task will be decremented.

- ii. (Rule – Timely boost) All processes in the system will be moved to the highest priority level periodically.

4. [Adapted from AY1920S1 Midterm – Evaluating scheduling algorithms]

Briefly answer each of the following questions regarding process scheduling, stating your assumptions, if any.

- a. Under what conditions does round-robin (RR) scheduling behave identically to FIFO?
- b. Under what conditions does RR scheduling perform poorly compared to FIFO?
- c. Under what conditions does FCFS (FIFO) scheduling result in the shortest possible average response time?

Questions for your own exploration (will not be discussed in the tutorial)

- 5. [Putting it together] Take a look at the given mysterious program **Behavior.c**. This program takes in one integer command line argument **D**, which is used as a **delay** to control the amount of computation work done in the program. For the part (a) and (b), use ideas you have learned from **Lecture 4: Process Scheduling** to explain the program behavior.

Use the command `taskset --cpu-list 0 ./Behaviors D`

This restricts the process to run on only one core.

Warning: you may not have the `taskset` command on your Linux system. If so, install the `util-linux` package using your package manager (`apt`, `yum`, etc).

Note: If you are using Windows Subsystem for Linux (WSL), make sure that you are using WSL2 kernel instead of WSL1. You can check by running `wsl -l -v` and upgrade using `wsl --set-version <distro-name> 2`

- a. **D** = 1.
- b. **D** = 100,000,000 (note: don't type in the ", " 😊)
- c. Now, find the **smallest D** that gives you the following interleaving output pattern:

Interleaving Output Pattern
[6183]: Step 0
[6184]: Step 0
[6183]: Step 1
[6184]: Step 1

```
[6183]: Step 2
[6184]: Step 2
[6183]: Step 3
[6184]: Step 3
[6183]: Step 4
[6184]: Step 4
[6184] Child Done!
[6183] Parent Done!
```

What do you think "**D**" represents?

*Note: "**D**" is machine dependent, you may get very different value from your friends'.*